

Traffic Sign Detection System

CSCI 331 - Group 2
RIT - 12/3/25

Miko Miguel D, Bishop O, Jonah C

Learning Goals and Task

Learning goals for this project:

- Understanding how machine learning models detect objects in real-world images (main focus on the traffic signs)
- Accurately classifying traffic signs under real-world conditions using robust and reliable perception modules.
- Learning the **ML pipeline**:
 - Preprocessing -> training -> evaluation


Assigned task:

- Build a system that detects and classifies multiple types of traffic signs from images using at least two AI/ML algorithm
- Compare the algorithms that we used for this project
- Developing a high fidelity instance segmentation solution for traffic signs

Background

Over the past few years, we've had a significant shift in computer vision techniques for object recognition. With many modern systems now utilizing Deep Convolutional Neural Networks (DCNNs)

- **Object Detection:** These focus on predicting bounding boxes and class labels. (Fast R-CNN)
- **Semantic Segmentation:** Classifying each pixels in an image, while treating all instances of the same object as a whole region. (U-net)
- **Instance Segmentation:** A combination of both the object detection and semantic segmentation model. Since it detects individual objects ad provides a unique mask for each instance. (Mask R-CNN)



Datasets Used for this Project

The Roboflow Traffic sign Dataset:

- Contains thousands of labelled traffic sign images from real driving photos
- Included **classes**:
 - Speed limit (10 - 120)
 - Stop signs
 - Traffic lights
- Each image comes with YOLO formatted bounding box annotations
- Dataset prematurely split for us:
 - Training
 - Validation
 - Test

1	7	0.5192307692307693	0.5384615384615384	0.5612980769230769	0.5540865384615384
2	6	0.5132211538461539	0.5360576923076923	0.578125	0.5576923076923077

The Method

What we tested:

Evaluation of different computer vision models' ability to detect and classify various traffic signs, comparing their final performance metrics against each other for the following:

- YoloV8
- Mask R-CNN

Our initial Hypothesis: Based on what we've seen of the models online, we expect the Mask R-CNN model to generally perform better than the YoloV8 model.

Metrics for evaluation

- Primary metrics:

- Precision Metric: True Positive / True Positive + False Positive
- Recall Metric: True Positive / True Positive + False Negative
- F1 score Metric: The harmonic mean of the Precision and Recall values

Accuracy	Predictions/ Classifications	$\frac{\text{Correct}}{\text{Correct} + \text{Incorrect}}$
Precision	Predictions/ Classifications	$\frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$
Recall	Predictions/ Classifications	$\frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$
F1	Predictions/ Classifications	$\frac{2 * \text{True Positive}}{\text{True Positive} + 0.5 (\text{False Positive} + \text{False Negative})}$
IoU	Object Detections/ Segmentations	$\frac{\text{Pixel Overlap}}{\text{Pixel Union}}$ 

Tools and Environments (YoloV8)

- Core Framework: PyTorch (used internally by Ultralytics)
- Vision Library: Ultralytics YOLOv8
- Monitoring: Built-in YOLO training logs
- Evaluation: mAP, Precision, Recall
- Hardware: CUDA GPU acceleration

Algorithm and Architecture (YoloV8)

- **Algorithm:** Object detection using Yolov8-n architecture
- **Base Model:** YOL0v8-n (nano model, fast, lightweight)
- **Pre-Training:** yes (yolov8n.pt)
- **Head Customization:** YOL0v8 automatically adapts output classes
- **Hyperparameters/Training**
 - Epochs: 20
 - Batch size: 16
 - Learning rate: default YOLO LR
 - Input size: default 640×640



Visualization (pt. 1: YoloV8)

- **Findings:**

- Overall strong performance
 - Precision of 0.94
 - Recall of 0.84
 - IoU: ~0.5 AND mAP50: 0.93
 - F1 = 0.89
- High accuracy for most classes with very few miscalculations
- PR curve has excellent stability
- Can detect signs in real images

- **Conclusions:**

- YoloV8 proves to be fast, accurate, and effective for how long I trained for (20 epochs)
- Shows strong generalization



Tools and Environments (Mask R-CNN)

- Core Framework: Pytorch (torch)
- Vision Library: torchvision (for model and transforms)
- Monitoring: tqdm (recording batch-level training progress)
- Evaluation: numpy, sklearn.metrics
- Hardware: CUDA (detects for GPU), fall back to CPU (torch.device)

Algorithm and Architecture (Mask R-CNN)

- **Algorithm:** Instance Segmentation using Mask R-CNN
- **Base Model:** Mask R-CNN with ResNet-50-FPN backbone (`maskrcnn_resnet50_fpn`)
- **Pre-Training:** Loaded model with `weight = None`
- **Head Customization:** Replaced both the Box Head (Fast RCNN Predictor) and Mask Head (Mask RCNN Predictor) to output $N = 2$ classes
- **Hyperparameters/Training**
 - `Num_Classes: 2` (Background and traffic sign)
 - `Optimizer: Stochastic Gradient Descent`
 - `Learning Rate: 0.005`
 - `Momentum: 0.9`
 - `Weight Decay: 0.0005`
 - `Num Epochs: 5`
 - `Batch Size: 2(training), 1(validation)`


Visualization (pt. 2: Mask R-CNN)

- **Findings:**


- Overall strong performance
 - Precision of 0.920
 - Recall of 0.880
 - F1 = 0.900
 - Training Loss of 0.00
- Low loss indicates a successful object localization and mask prediction.
- Model rarely misidentified background as a traffic sign due to its high precision value.
- Successfully found most traffic signs present in each image due to its high recall rate.

- **Conclusions:**

- Mask R-CNN proved a successful convergence over 5 training epochs.
- Real-time tracking proved evident that the network was actively learning and not hanging. (More accurate at detecting and segmenting traffic signs)
- x-axis(epoch) # of times the model has seen the entire training dataset while y-axis (Loss) amount of error made by the model.



Sample Output Images from Validation Set



Takeaways and insight

- Our Strengths:
 - Mask R-CNN is better suited for identifying and separating multiple distinct traffic signs in a single image
 - Yolo offers higher precision at the slight but equal cost of recall
 - Time taken to train was much faster as we had the time to train 20 epochs in Yolo in the time we took to train 5 epochs in Mask R-CNN
- Limitations we noticed:
 - Yolo can't easily separate or distinguish between two instances.
 - Image-level Precision/Recall/F1 only measures whether the model made a detection when a sign was present, it doesn't fully measure the accuracy of the predicted bounding box or mask.
 - Smaller Batch Size can lead to noisy gradient estimates
 - The efficiency of training process is too dependent on the size and diversity of the Kaggle data set. This could lead to poor generalization beyond training samples.

Conclusion - Discussion

Our big takeaways:

- Transitioning from pixel-level ground truth to structured object tensors like boxes, masks, and labels helped in achieving a high F1 score for the binary detection.

Stuff to look into if we had more time:

- Mean Average Precision and IoU calculation for instance segmentation

Acknowledgement

- Darabi, Parisa Karimi. “Traffic Signs Detection.” *Kaggle.com*, 2024, www.kaggle.com/datasets/pkdarabi/cardetection/data.
- “Torchvision – Torchvision 0.24 Documentation.” *Pytorch.org*, 2024, docs.pytorch.org/vision/stable/. Accessed 3 Dec. 2025.
- Mohammed, Samiyaa Yaseen. “Architecture Review: Two-Stage and One-Stage Object Detection.” *Franklin Open*, vol. 12, 17 July 2025, p. 100322, www.sciencedirect.com/science/article/pii/S2773186325001100, <https://doi.org/10.1016/j.fraope.2025.100322>.
- *Mask R-CNN*. (2017, October 1). IEEE Conference Publication | IEEE Xplore.
<https://ieeexplore.ieee.org/document/8237584>

Any Questions?